

Adwait Minde

+1 (520) 283 8617 • adwaitm131@gmail.com • www.linkedin.com/in/adwait-minde/

EDUCATION

Master of Science in Management Information Systems
Bachelor of Engineering in Electronics and Computer Science

University of Arizona
University of Mumbai

GPA: 3.89
GPA: 3.47

Dec 2025
May 2023

SKILLS

Programming & Scripting: Python, PostgreSQL, Oracle SQL, C, C++, JavaScript, HTML, CSS, React, FastAPI

Data Analysis & Machine Learning: Pandas, NumPy, Scikit-learn, PyTorch, Hugging Face, NLP, Data Mining

Databases & ETL: Oracle SQL, Redshift, Apache Airflow, AWS Glue, Athena, Data Pipeline Development, ETL Automation

Visualization & Reporting: Power BI, Tableau, MS Excel (Pivot Tables, VLOOKUP, Data Modeling), MS Word, Microsoft Access

WORK EXPERIENCE

Data Analyst | Banner Health | Python, Pandas, Plotly, LLM-Powered Insights

Jun 2025 - Aug 2025

- Analyzed chronic disease prevalence across states using CMS and census datasets to identify health disparities by age, race, and county
- Built Python workflows with Pandas to visualize condition trends (COPD, Hypertension) via choropleths and stratified bar charts
- Validated outputs against CDC benchmarks to assess model consistency, applying algorithm evaluation techniques and statistical checks
- Delivered interactive visual dashboards with Plotly and Seaborn to support data-driven decisions on healthcare outreach and planning

Team Lead | Arizona Public Services | Excel, Power BI, Financial Modeling

Jan 2025 – May 2025

- Led a 6-member team to evaluate 4 capital projects totaling \$250M, shaping long-term strategy for a nuclear facility
- Managed timelines and deliverables using MS Project and JIRA, ensuring on-time execution and streamlined coordination
- Built financial models in Excel to project ROI and NPV, identifying \$1.2M+ in potential annual savings
- Created Power BI dashboards to visualize cost, risk, and utilization metrics, enabling informed executive decisions

Business Development Analyst | Zycus Pvt. Ltd.

Jun 2023 – May 2024

- Built batch ingestion pipelines leveraging ERP APIs to integrate POs, invoices, contracts, and supplier master data into a centralized data lake, enabling scalable processing and analysis across 180+ retail clients via Salesforce integration.
- Configured Apache Kafka consumers to process procurement event streams, reducing data latency for operational reporting
- Applied PySpark transformations to standardize raw spend data formats across, cutting manual data prep effort by approximately 20%
- Implemented dbt models alongside analytics engineers to structure raw procurement data into reusable dimension and fact tables, reducing ad-hoc query development time across the BI team
- Wrote SQL transformations for spend categorization logic, improving classification accuracy in the AI-powered Spend Analysis module
- Developed supplier deduplication logic using SQL window functions, eliminating an estimated 18% duplicate supplier record
- Leveraged AWS S3, Redshift, and Glue to build a scalable ETL infrastructure, supporting analytical queries on procurement spend data

ACADEMIC EXPERIENCE

Project Lead | NYC Taxi Lakehouse: End-to-End Medallion Pipeline | PySpark, Delta Lake, Databricks Jobs

May 2026 – Jun 2026

- Engineered a bronze/silver/gold medallion pipeline processing 270M+ NYC TLC trip records across 4 transportation datasets, eliminating downstream type-conversion failures by validating Delta schemas against production files rather than vendor documentation
- Designed a deterministic data quality layer with PySpark window-function deduplication and quarantine routing, then root-caused a Spark casting defect that had been silently corrupting trip-duration calculations across 5 pipeline stages
- Orchestrated a 12-task dependency-aware DAG in Databricks Jobs, structuring task dependencies to let independent aggregates run without waiting on the slowest upstream stage, cutting unnecessary pipeline latency
- Deployed a public-facing Streamlit dashboard on the gold layer, giving stakeholders live, self-service access to revenue, demand, and trip-duration insights without needing direct warehouse access; live at nyc-taxi-lakehouse-dashboard.streamlit.app

Project Lead | Real-Time Customer Churn Intelligence Platform | Kafka, BigQuery, dbt, XGBoost, Claude AI, Streamlit

Jun 2026

- Built an end-to-end streaming pipeline using Kafka and AWS S3 to ingest 500K e-commerce behavioral events, with a dbt-powered BigQuery warehouse producing a labeled customer feature mart across 7,362 users with 7 schema tests enforcing data quality
- Trained an XGBoost churn classifier across 3 MLflow-tracked experiments, achieving AUC 0.9956 and F1 0.9911, with SHAP analysis identifying recency and purchase frequency as the dominant churn signals
- Integrated Claude AI (Haiku) to generate per-customer plain-English churn explanations for 6,826 at-risk users, writing outputs back to BigQuery to make model predictions directly actionable for non-technical retention teams
- Deployed a two-tab Streamlit dashboard on Streamlit Cloud with interactive Plotly charts, live BigQuery integration, and risk-tiered customer explorer; accessible at churn-intelligence-platform-project.streamlit.app

Project Lead | Data Pipelines with Airflow | Apache Airflow, Amazon Redshift, S3, Python

Dec 2024 - Jan 2025

- Built automated ETL pipelines in Airflow with custom operators to load S3 data into Redshift, enabling daily automated refreshes
- Improved reliability by designing idempotent, parameterized workflows with partitioned loads, cutting reprocessing time by 40%
- Strengthened data integrity with a DataQualityOperator that ran SQL checks, catching anomalies before reaching dashboards
- Delivered a star-schema warehouse integrated with BI tools (Tableau, QuickSight), supporting timely and accurate stakeholder insights

Project Head | Data management of McDonald's supply chain | SQL, Oracle APEX

Aug 2024 - Dec 2024

- Analyzed inefficiencies in the supply chain data system and designed an ER model using MS Visio to optimize data relationships
- Converted the ER model into a relational schema, created normalized tables, and applied 4NF to ensure data consistency
- Developed SQL triggers and stored procedures to automate data input, reducing manual effort and ensuring data integrity